

Time statistics

IMPROVING QUERY PERFORMANCE IN SQL SERVER



Dean Smith

Founder, Atamai Analytics

SQL Server Management Studio

Microsoft
SQL Server Management Studio

STATISTICS TIME in SSMS

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital  
FROM Nations  
WHERE Capital IN  
       (SELECT CityName  
        FROM Cities)  
       AND Capital IN  
        (SELECT CityName  
         FROM Cities  
         WHERE Pop2017 > 1000000);
```

STATISTICS TIME - command reports number of milliseconds, required to parse, compile, and execute a query.

SQL Server parse and compile time:

CPU time = 16 ms, elapsed time = 21 ms.

(88 rows affected)

SQL Server Execution Times:

CPU time = 390 ms, elapsed time = 382 ms.

SQL Server Execution Times

SQL Server Execution Times:

CPU time = 390 ms, elapsed time = 382 ms.

- **CPU time:** time taken by server processors to process the query
- **Elapsed time:** total duration of the query

Example: query 1

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital  
FROM Nations  
WHERE Capital IN  
      (SELECT CityName -- 1st sub-query  
       FROM Cities)  
      AND Capital IN  
      (SELECT CityName -- 2nd sub-query  
       FROM Cities  
       WHERE Pop2017 > 1000000);
```

Example: query 1

```
SET STATISTICS TIME ON
```

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital  
FROM Nations  
WHERE Capital IN  
      (SELECT CityName -- 1st sub-query  
       FROM Cities)  
      AND Capital IN  
      (SELECT CityName -- 2nd sub-query  
       FROM Cities  
       WHERE Pop2017 > 1000000);
```

SQL Server Execution Times:

CPU time = 391 ms, elapsed time = 381 ms.

Example: query 2

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital  
FROM Nations n  
WHERE EXISTS  
  (SELECT 1  
   FROM Cities c  
   WHERE n.Capital = c.CityName  
        AND Pop2017 > 1000000);
```

SQL Server Execution Times:

CPU time = 0 ms, elapsed time = 2 ms.

Example: query 2

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital  
FROM Nations n  
WHERE EXISTS  
  (SELECT 1  
   FROM Cities c  
   WHERE n.Capital = c.CityName  
        AND Pop2017 > 1000000);
```

```
SET STATISTICS TIME OFF
```

SQL Server Execution Times:

CPU time = 0 ms, elapsed time = 2 ms.

Comparing queries

- First query containing two sub-queries

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital  
FROM Nations  
WHERE Capital IN  
      (SELECT CityName  
       FROM Cities)  
      AND Capital IN  
      (SELECT CityName  
       FROM Cities  
       WHERE Pop2017 > 1000000);
```

SQL Server Execution Times:

CPU time = 391 ms, elapsed time = 381 ms.

- Second query using EXISTS

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital  
FROM Nations n  
WHERE EXISTS  
      (SELECT 1  
       FROM Cities c  
       WHERE n.Capital = c.CityName  
              AND Pop2017 > 1000000);
```

SQL Server Execution Times:

CPU time = 0 ms, elapsed time = 2 ms.

Elapsed time vs. CPU time

Elapsed time

- May be variable when analyzing query time statistics
- Time best time statistics measure for the fastest running query

CPU time

- Should vary little when analyzing query time statistics
- May not be a useful measure if server processors are running in parallel

Taking an average

- Don't rely on one measurement, take an average.

```
elapsed time = 2032 ms.
```

```
elapsed time = 2060 ms.
```

```
elapsed time = 1915 ms.
```

```
elapsed time = 4009 ms.
```

```
elapsed time = 3511 ms.
```

```
Average elapsed time = 2705 ms.
```

Let's practice!

IMPROVING QUERY PERFORMANCE IN SQL SERVER

Page read statistics

IMPROVING QUERY PERFORMANCE IN SQL SERVER



Dean Smith

Founder, Atamai Analytics

Table data pages

- All data, in either memory or on the disk, is stored in 8 kilobyte size "pages"
- One page can store many rows or one value could span multiple pages
- A page can only belong to one table
- SQL Server works with pages cached in memory
- If a page is not cached in memory it is read from disk and cached in memory

Customers: data pages

Page 1

```
| ALFKI | Alfreds Futterkiste | Maria Anders | Sales Representative | ...  
| ANATR | Ana Trujillo Emparedados y helados | Ana Trujillo | Owner | ...  
| ANTON | Antonio Moreno Taqueria | Antonio Moreno | Owner | ...  
| AROUT | Around the Horn | Thomas Hardy | Sales Representative | ...  
| BERGS | Berglunds snabbkop | Christina Berglund | Order Administrator | ...  
| BLAUS | Blauer See Delikatessen | Hanna Moos | Sales Representative | ...  
| BLONP | Blondesddsl pere et fils | Frederique Citeaux | Marketing Manager | ...
```

Page 2

```
| BOLID | Bolido Comidas preparadas | Martin Sommer | Owner | ...  
| BONAP | Bon app | Laurence Lebihan | Owner | ...  
| BOTTM | Bottom-Dollar Markets | Elizabeth Lincoln | Accounting Manager | ...  
| BSBEV | B s Beverages | Victoria Ashworth | Sales Representative | ...  
| CACTU | Cactus Comidas para llevar | Patricio Simpson | Sales Agent | ...  
| CENTC | Centro comercial Moctezuma | Francisco Chang | Marketing Manager | ...  
| CHOPS | Chop-suey Chinese | Yang Wang | Owner | ...
```

Page 3

```
| COMMI | Comercio Mineiro | Pedro Afonso | Sales Associate | ...  
| CONSH | Consolidated Holdings | Elizabeth Brown | Sales Representative | ...  
| DRACD | Drachenblut Delikatessen | Sven Ottlieb | Order Administrator | ...  
| DUMON | Du monde entier | Janine Labrune | Owner | ...  
| EASTC | Eastern Connection | Ann Devon | Sales Agent | ...  
| ERNSH | Ernst Handel | Roland Mendel | Sales Manager | ...  
| FAMIA | Familia Arquibaldo | Aria Cruz | Marketing Assistant | ...
```

Page 4

```
| FISSA | FISSA Fabrica Inter. Salchichas S.A. | Diego Roel | Accounting Manager | ...  
| FOLIG | Folies gourmandes | Martine Rance | Assistant Sales Agent | ...  
| FOLKO | Folk och fa HB | Maria Larsson | Owner | ...  
| FRANK | Frankenversand | Peter Franken | Marketing Manager | ...  
| FRANR | France restauration | Carine Schmitt | Marketing Manager | ...  
| FRANS | Franchi S.p.A. | Paolo Accorti | Sales Representative | ...  
| FURIB | Furia Bacalhau e Frutos do Mar | Lino Rodriguez | Sales Manager | ...
```

STATISTICS IO in SSMS

```
SET STATISTICS IO ON
```

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital  
FROM Nations  
WHERE Capital IN  
      (SELECT CityName  
       FROM Cities)  
      AND Capital IN  
      (SELECT CityName  
       FROM Cities  
       WHERE Pop2017 > 1000000);
```

(88 rows affected)

Table 'Worktable'. Scan count 1, logical reads 104568, physical reads 0, read-ahead reads 0, lob logical reads 0, ...

Table 'Cities'. Scan count 2, logical reads 548, physical reads 0, read-ahead reads 281, lob logical reads 0, lob physical reads 0, ...

Table 'Nations'. Scan count 1, logical reads 3, physical reads 1, read-ahead reads 0, lob logical reads 0, lob physical reads 0, ...

Logical reads

(88 rows affected)

Table 'Cities'. ... , logical reads 548, ... , ...

Table 'Nations'. ... , logical reads 3, ... , ...

logical reads: number of 8 kilobyte pages read per table

Example: query 1

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital,  
       (SELECT MAX(Magnitude)  
        FROM Earthquakes e  
        WHERE n.Capital = e.NearestPop  
              AND n.Code2 = e.Country) MaxMagnitude  
FROM Nations n  
WHERE Capital IN  
      (SELECT NearestPop  
       FROM Earthquakes e  
       WHERE n.Capital = e.NearestPop  
             AND n.Code2 = e.Country)  
      AND Capital IN  
            (SELECT NearestPop  
             FROM Earthquakes  
             WHERE Magnitude >= 7.5);
```

Example: query 1

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital,  
       (SELECT MAX(Magnitude) -- 1st sub-query  
        FROM Earthquakes e  
        WHERE n.Capital = e.NearestPop  
              AND n.Code2 = e.Country) MaxMagnitude  
FROM Nations n  
WHERE Capital IN  
      (SELECT NearestPop  
       FROM Earthquakes e  
       WHERE n.Capital = e.NearestPop  
            AND n.Code2 = e.Country)  
      AND Capital IN  
        (SELECT NearestPop  
         FROM Earthquakes  
         WHERE Magnitude >= 7.5);
```

Example: query 1

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital,  
       (SELECT MAX(Magnitude) -- 1st sub-query  
        FROM Earthquakes e  
        WHERE n.Capital = e.NearestPop  
              AND n.Code2 = e.Country) MaxMagnitude  
FROM Nations n  
WHERE Capital IN  
      (SELECT NearestPop -- 2nd sub-query  
       FROM Earthquakes e  
       WHERE n.Capital = e.NearestPop  
             AND n.Code2 = e.Country)  
      AND Capital IN  
            (SELECT NearestPop  
             FROM Earthquakes  
             WHERE Magnitude >= 7.5);
```

Example: query 1

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital,  
       (SELECT MAX(Magnitude) -- 1st sub-query  
        FROM Earthquakes e  
        WHERE n.Capital = e.NearestPop  
              AND n.Code2 = e.Country) MaxMagnitude  
FROM Nations n  
WHERE Capital IN  
      (SELECT NearestPop -- 2nd sub-query  
       FROM Earthquakes e  
       WHERE n.Capital = e.NearestPop  
             AND n.Code2 = e.Country)  
      AND Capital IN  
            (SELECT NearestPop -- 3rd sub-query  
             FROM Earthquakes  
             WHERE Magnitude >= 7.5);
```

Example: query 1

```
SET STATISTICS IO ON
```

```
SELECT UNStatisticalRegion,  
       CountryName,  
       Capital,  
       (SELECT MAX(Magnitude) -- 1st sub-query  
        FROM Earthquakes e  
        WHERE n.Capital = e.NearestPop  
              AND n.Code2 = e.Country) MaxMagnitude  
FROM Nations n  
WHERE Capital IN  
      (SELECT NearestPop -- 2nd sub-query  
       FROM Earthquakes e  
       WHERE n.Capital = e.NearestPop  
            AND n.Code2 = e.Country)  
      AND Capital IN  
      (SELECT NearestPop -- 3rd sub-query  
       FROM Earthquakes  
       WHERE Magnitude >= 7.5);
```

Table 'Earthquakes'. ..., logical reads 54, ...
'Nations'. ..., logical reads 3, ...

Example: query 2

```
SELECT n.UNStatisticalRegion,  
       n.CountryName,  
       n.Capital,  
       MAX(e.Magnitude)  
FROM Nations n  
INNER JOIN Earthquakes e  
    ON n.Capital = e.NearestPop  
    AND n.Code2 = e.Country  
WHERE e.Magnitude >=7.5  
GROUP BY n.UNStatisticalRegion,  
         n.CountryName,  
         n.Capital
```

Example: query 2

```
SELECT n.UNStatisticalRegion,  
       n.CountryName,  
       n.Capital,  
       MAX(e.Magnitude)  
FROM Nations n  
INNER JOIN Earthquakes e  
  ON n.Capital = e.NearestPop  
   AND n.Code2 = e.Country  
WHERE e.Magnitude >=7.5  
GROUP BY n.UNStatisticalRegion,  
         n.CountryName,  
         n.Capital
```

```
SET STATISTICS IO OFF
```

```
Table 'Earthquakes'. ..., logical reads 18, ...  
'Nations'. ..., logical reads 3, ...
```


Comparing queries

- First query containing three sub-queries

```
Table 'Earthquakes'. ..., logical reads 54, ...  
'Nations'. ..., logical reads 3, ...
```

SQL Server parse and compile time:

CPU time = 29 ms, elapsed time = 29 ms.

- Second query using an `INNER JOIN`

```
Table 'Earthquakes'. ..., logical reads 18, ...  
'Nations'. ..., logical reads 3, ...
```

SQL Server Execution Times:

CPU time = 0 ms, elapsed time = 3 ms.

Let's practice!

IMPROVING QUERY PERFORMANCE IN SQL SERVER

Indexes

IMPROVING QUERY PERFORMANCE IN SQL SERVER



Dean Smith

Founder, Atamai Analytics

What is an index?

- Structure to improve speed of accessing data from a table
- Used to locate data quickly without having to scan the entire table
- Useful for improving performance of queries with filter conditions
- Applied to table columns
- Typically added by a database administrator

Clustered and nonclustered indexes

Clustered Index

- Analogy : dictionary
- Table data pages are ordered by the column(s) with the index
- Only one allowed per table
- Speeds up search operations

Clustered and nonclustered indexes

Clustered Index

- Analogy : dictionary
- Table data pages are ordered by the column(s) with the index
- Only one allowed per table
- Speeds up search operations

Non-clustered Index

- Analogy: text book with an index at the back
- Structure contains an ordered layer of index pointers to unordered table data pages
- A table can have more than one
- Improves insert and update operations

Clustered index: B-tree structure

- ROOT NODE
- BRANCH NODES
- PAGE NODES

Clustered index: B-tree structure

ROOT NODE:

A G O W

BRANCH NODES:

A B E F G H J K O P S T

PAGE NODES:

Page 1	Page 2	Page 3	Page 4
Index Column ...	Index Column ...	Index Column ...	Index Column ...
A ...	E ...	I ...	M ...
B ...	F ...	J ...	N ...
C ...	G ...	K ...	O ...
D ...	H ...	L ...	P ...
...

Customers table without clustered index

```
SELECT *  
FROM Customers  
WHERE CustomerID = "PARIS"
```

PAGE NODES:

```
    Page 1  
CustomerID | ...  
ALFKI    | ...  
ANATR    | ...  
BLONP    | ...  
BSBEV    | ...  
...
```

Customers table without clustered index

```
SELECT *  
FROM Customers  
WHERE CustomerID = "PARIS"
```

PAGE NODES:

Page 1	Page 2
CustomerID ...	CustomerID ...
ALFKI ...	FOLIG ...
ANATR ...	FRANK ...
BLONP ...	GALED ...
BSBEV ...	GREAL ...
...	...

Customers table without clustered index

```
SELECT *  
FROM Customers  
WHERE CustomerID = "PARIS"
```

PAGE NODES:

Page 1	Page 2	Page 3
CustomerID ...	CustomerID ...	CustomerID ...
ALFKI ...	FOLIG ...	LILAS ...
ANATR ...	FRANK ...	LINOD ...
BLONP ...	GALED ...	MEREP ...
BSBEV ...	GREAL ...	MORGK ...
...

Customers table without clustered index

```
SELECT *  
FROM Customers  
WHERE CustomerID = "PARIS"
```

PAGE NODES:

Page 1	Page 2	Page 3	Page 4
CustomerID ...	CustomerID ...	CustomerID ...	CustomerID ...
ALFKI ...	FOLIG ...	LILAS ...	OCEAN ...
ANATR ...	FRANK ...	LINOD ...	PARIS ...
BLONP ...	GALED ...	MEREP ...	
BSBEV ...	GREAL ...	MORGK ...	
...

Customers table with clustered index

```
SELECT *
FROM Customers
WHERE CustomerID = "PARIS"
```

ROOT NODE:

ALFKI FOLIG OLDWO WOLZA

BRANCH NODES:

ALFKI BONAP DRACD FISSA FOLIG GALED LILAS NORTS OCEAN OLDWO QUICK WOLZA

PAGE NODES:

Page 1	Page 2	Page 3	Page 4
CustomerID ...	CustomerID ...	CustomerID ...	CustomerID ...
ALFKI ...	FOLIG ...	LILAS ...	OCEAN ...
ANATR ...	FRANK ...	LINOD ...	PARIS ...
BLONP ...	GALED ...	MEREP ...	PICCO ...
BSBEV ...	GREAL ...	MORGK ...	QUICK ...
...

Customers table with clustered index

```
SELECT *
FROM Customers
WHERE CustomerID = "PARIS"
```

ROOT NODE:

OLDWO WOLZA

BRANCH NODES:

ALFKI BONAP DRACD FISSA FOLIG GALED LILAS NORTS OCEAN OLDWO QUICK WOLZA

PAGE NODES:

Page 1	Page 2	Page 3	Page 4
CustomerID ...	CustomerID ...	CustomerID ...	CustomerID ...
ALFKI ...	FOLIG ...	LILAS ...	OCEAN ...
ANATR ...	FRANK ...	LINOD ...	PARIS ...
BLONP ...	GALED ...	MEREP ...	PICCO ...
BSBEV ...	GREAL ...	MORGK ...	QUICK ...
...

Customers table with clustered index

```
SELECT *
FROM Customers
WHERE CustomerID = "PARIS"
```

ROOT NODE:

OLDWO WOLZA

BRANCH NODES:

OLDWO QUICK

PAGE NODES:

Page 1	Page 2	Page 3	Page 4
CustomerID ...	CustomerID ...	CustomerID ...	CustomerID ...
ALFKI ...	FOLIG ...	LILAS ...	OCEAN ...
ANATR ...	FRANK ...	LINOD ...	PARIS ...
BLONP ...	GALED ...	MEREP ...	PICCO ...
BSBEV ...	GREAL ...	MORGK ...	QUICK ...
...

Customers table with clustered index

```
SELECT *  
FROM Customers  
WHERE CustomerID = "PARIS"
```

ROOT NODE:

OLDWO WOLZA

BRANCH NODES:

OLDWO QUICK

PAGE NODES:

```
Page 4  
CustomerID | ...  
OCEAN | ...  
PARIS | ...  
PICCO | ...  
QUICK | ...  
...
```


Customers table with clustered index

```
SELECT *  
FROM Customers  
WHERE CustomerID = "PARIS"
```

ROOT NODE:

OLDWO WOLZA

BRANCH NODES:

OLDWO QUICK

PAGE NODES:

Page 4
CustomerID | ...

PARIS | ...

Clustered index: example

```
SET STATISTICS IO ON
```

```
SELECT *  
FROM PlayerStats  
WHERE Team = 'OKC'
```

PlayerStats table with no index

Table 'PlayerStats'. ..., logical reads 12, ...

PlayerStats table with **clustered index** on
Team

Table 'PlayerStats'. ..., logical reads 2, ...

Let's practice!

IMPROVING QUERY PERFORMANCE IN SQL SERVER

Execution plans

IMPROVING QUERY PERFORMANCE IN SQL SERVER



Dean Smith

Founder, Atamai Analytics

Optimization phase

Optimization Phase

Optimization phase

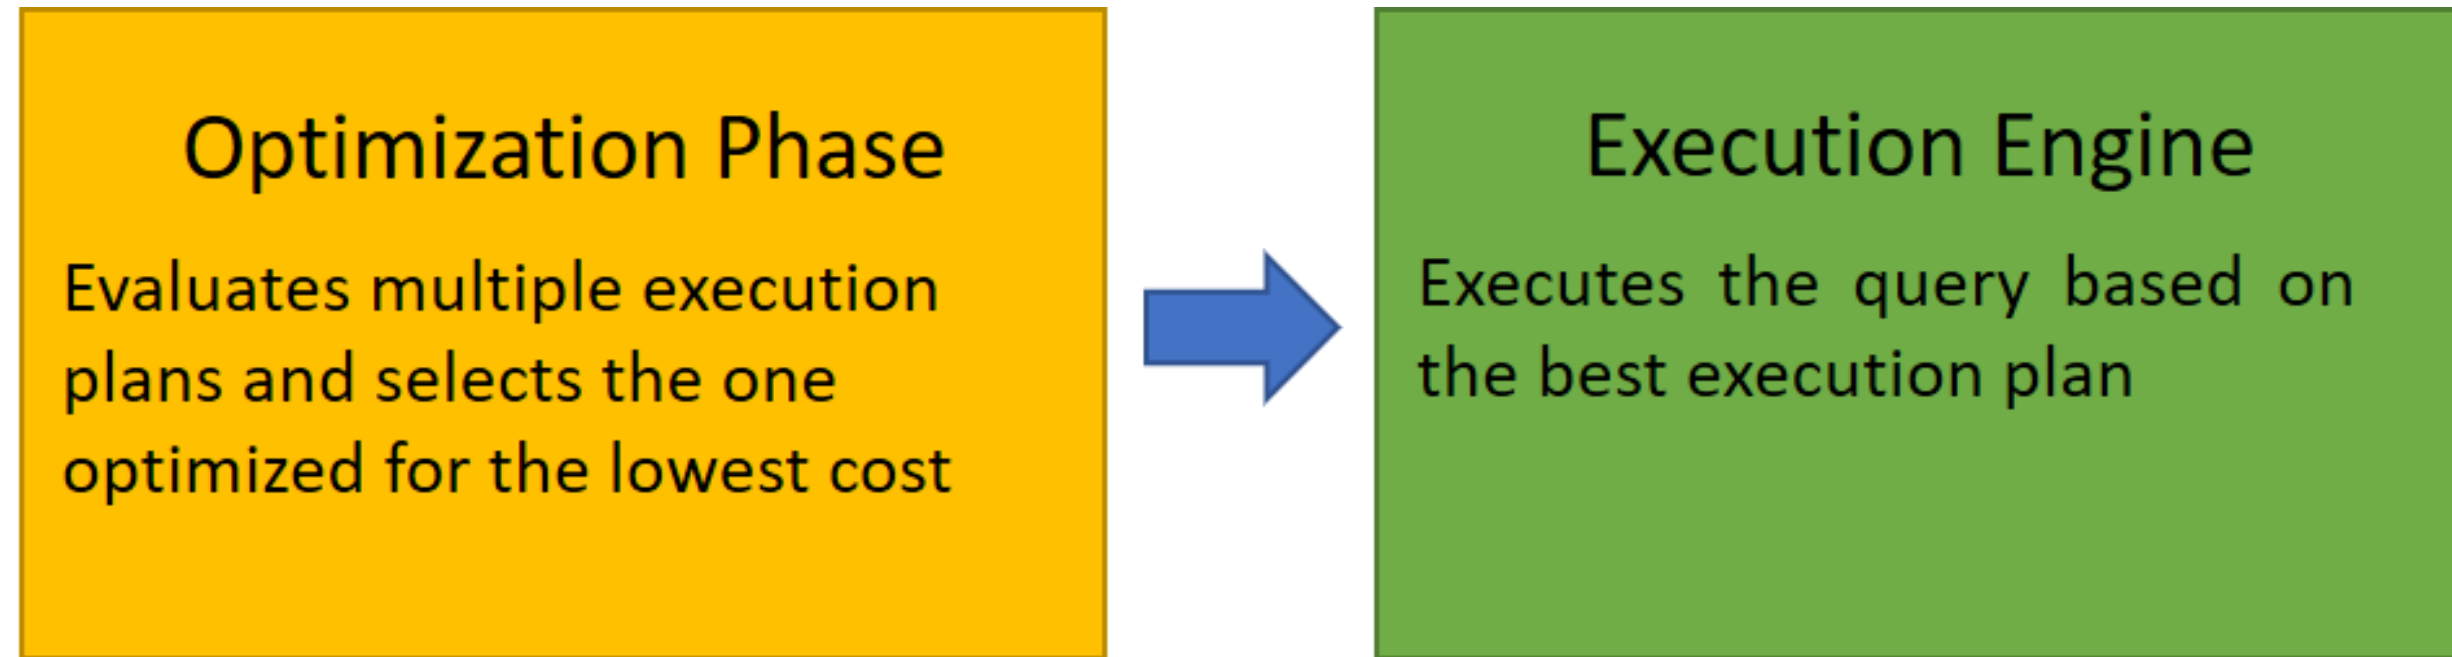
Optimization Phase

Evaluates multiple execution plans and selects the one optimized for the lowest cost

Cost parameters evaluated include:

- Processor usage
- Memory usage
- Data page reads

Optimization phase



Cost parameters evaluated include;

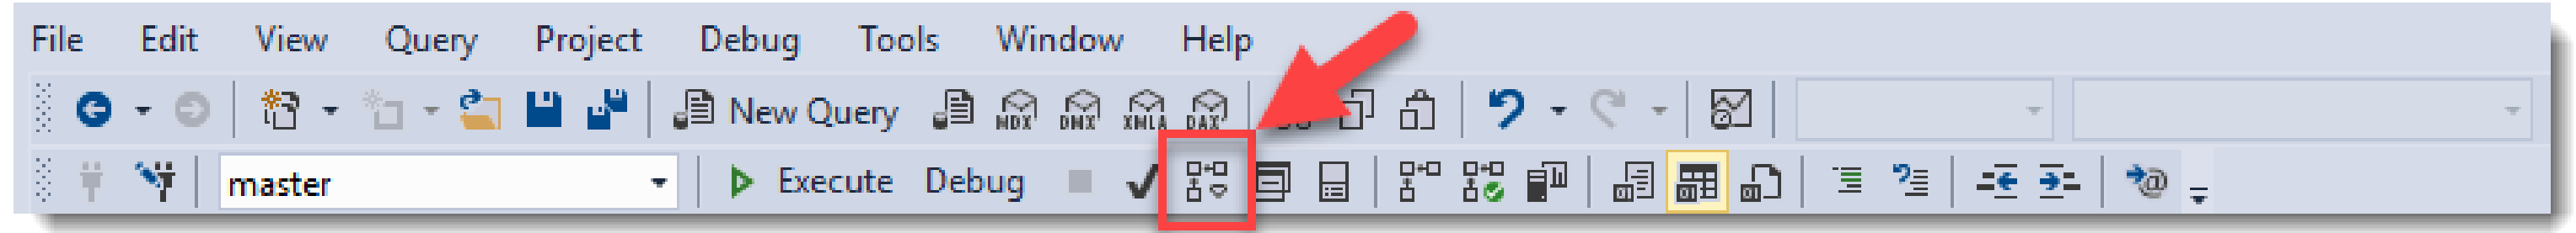
- Processor usage
- Memory usage
- Data page reads

Information from execution plans

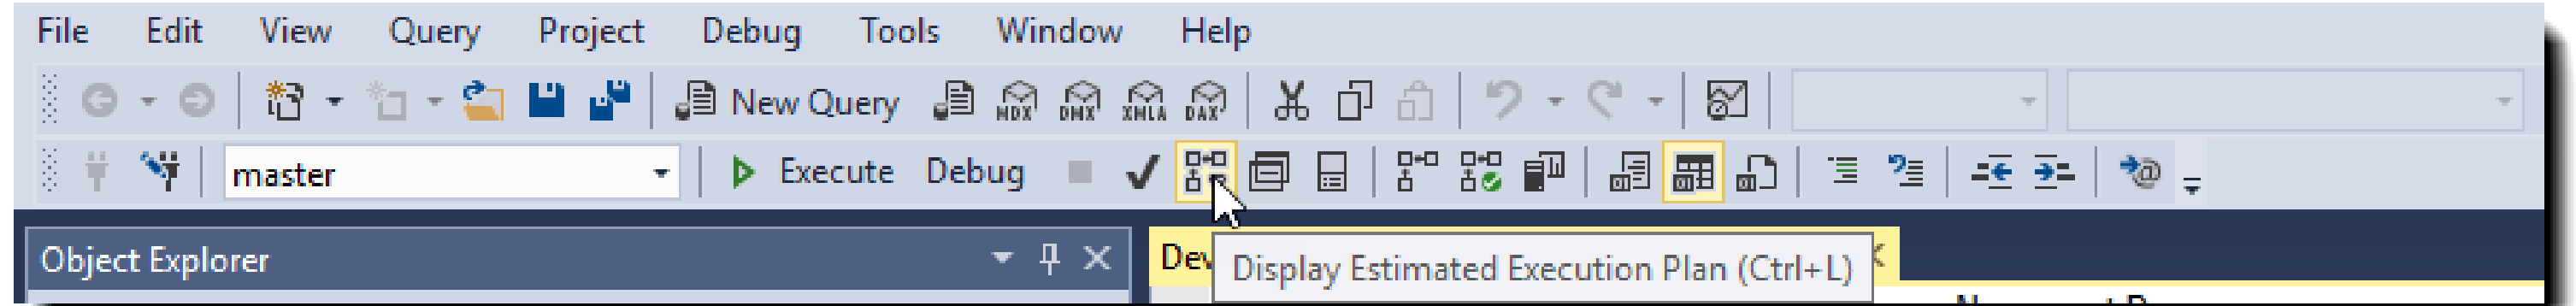
Execution plans can provide information on:

- Whether indexes were used
- Types of joins used
- Location and relative costs of:
 - filter conditions
 - sorting
 - aggregations

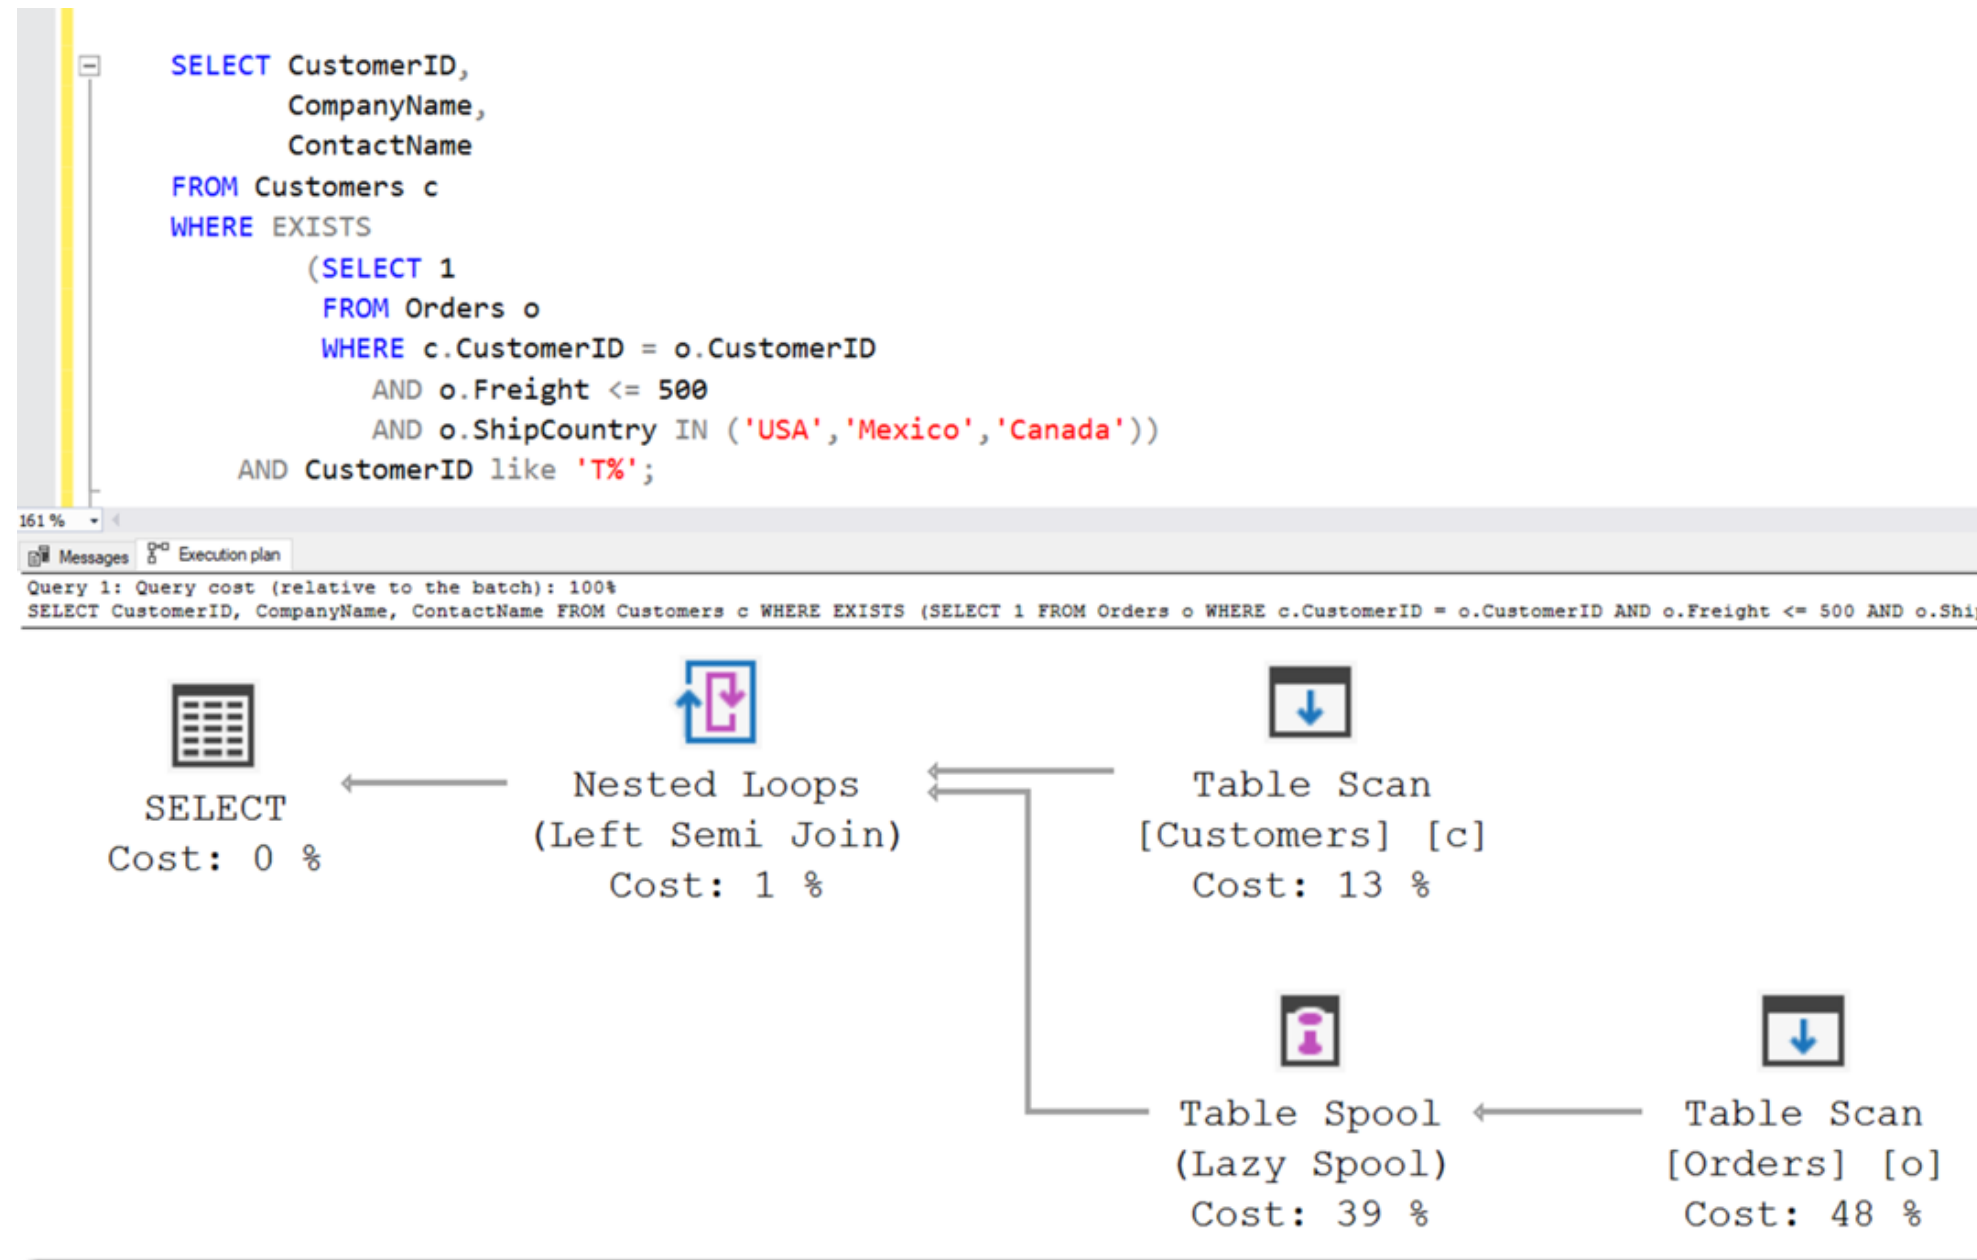
Estimated execution plan in SSMS



Estimated execution plan in SSMS



Viewing executions plans in SSMS



Operator statistics


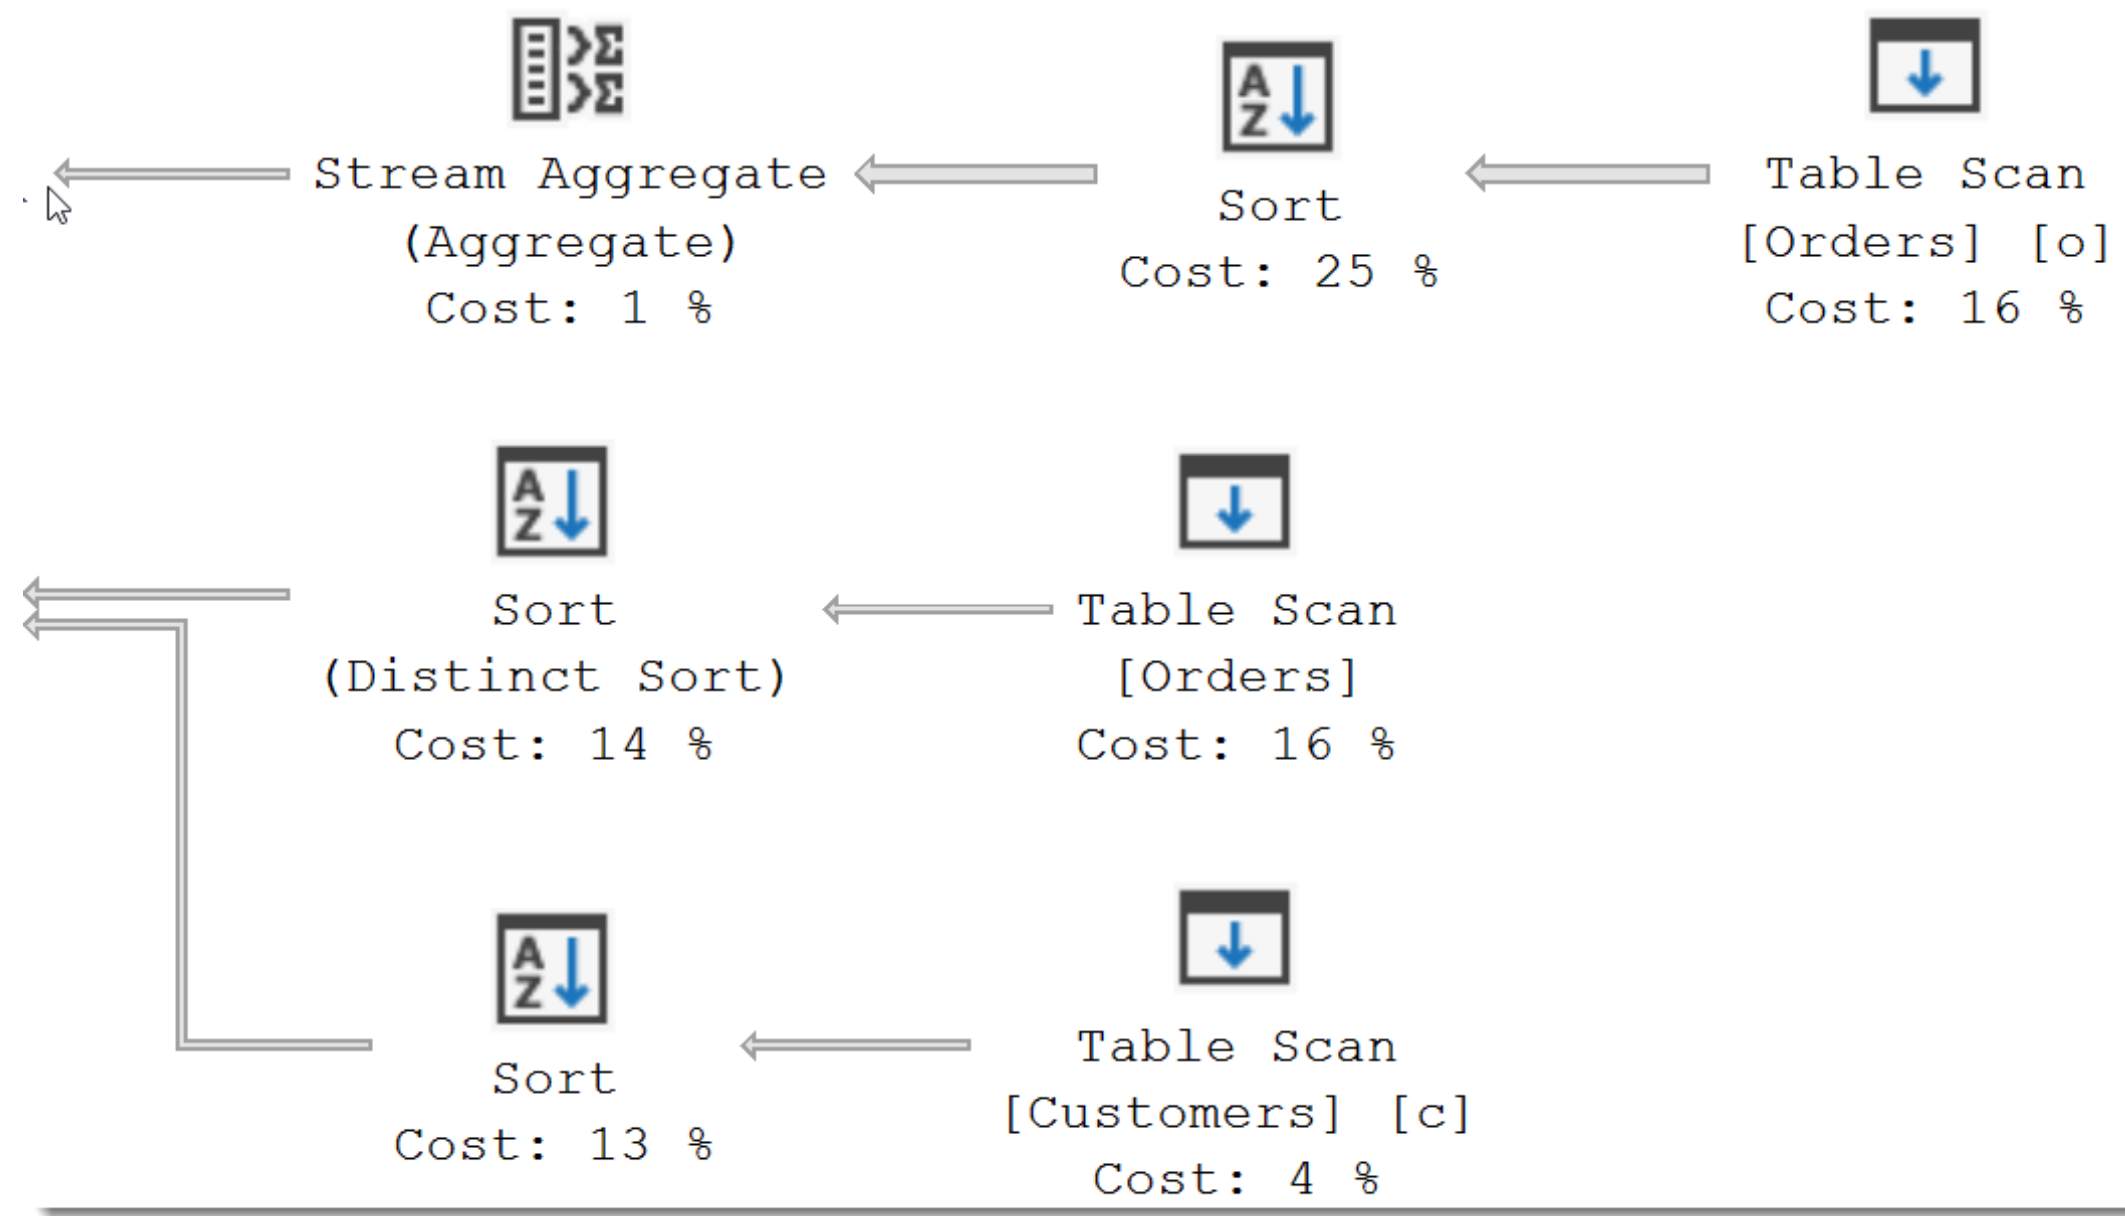



Table Scan	Table Scan
Scan rows from a table.	
Physical Operation	Table Scan
Logical Operation	Table Scan
Estimated Execution Mode	Row
Storage	RowStore
Estimated I/O Cost	0.0038657
Estimated Operator Cost	0.0041228 (13%)
Estimated CPU Cost	0.0002571
Estimated Subtree Cost	0.0041228
Estimated Number of Executions	1
Estimated Number of Rows	6.5
Estimated Number of Rows to be Read	91
Estimated Row Size	55 B
Ordered	False
Node ID	1

Reading execution plans

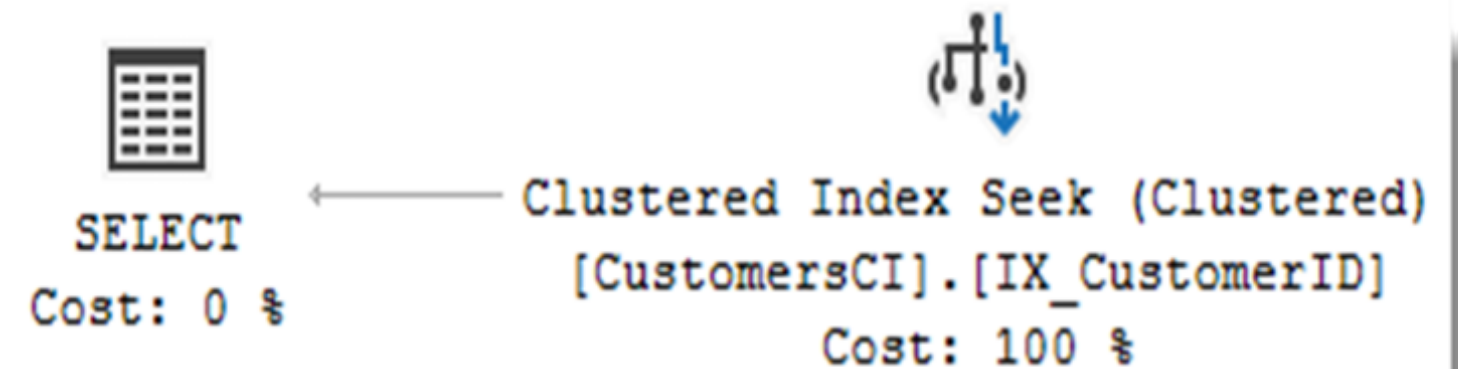


Index example

```
SELECT *  
FROM Customers  
WHERE CustomerID = 'PARIS';
```



```
SELECT *  
FROM CustomersCI  
WHERE CustomerID = 'PARIS';
```

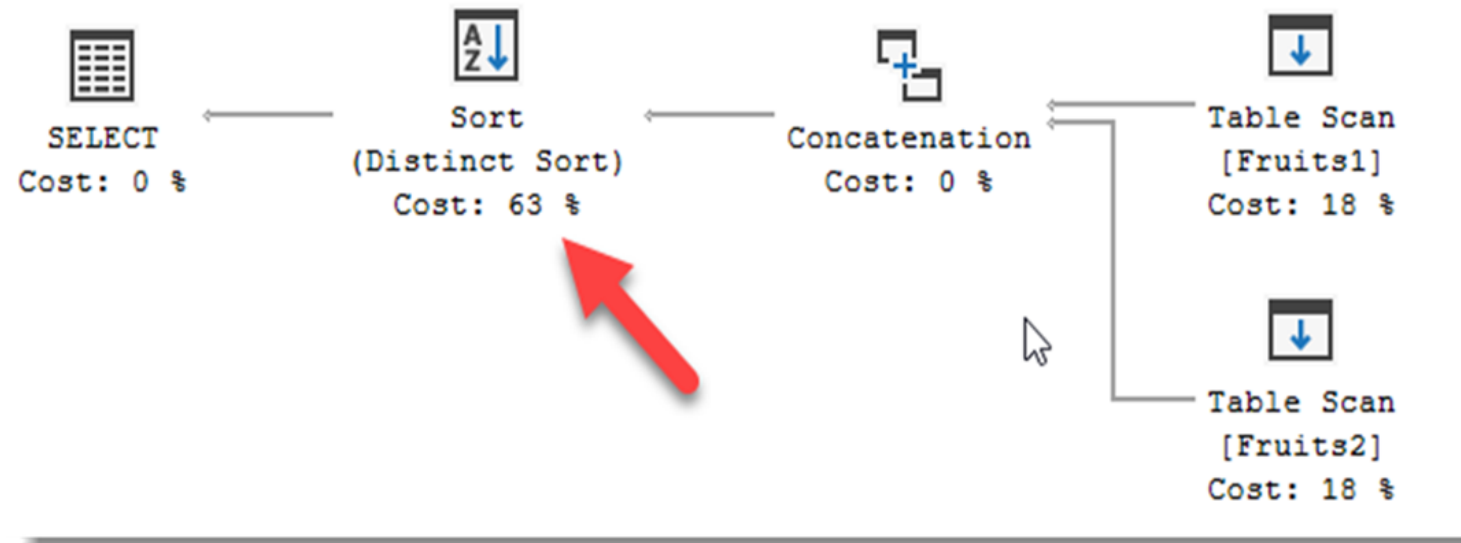


Sort operator example

```
SELECT FruitName, FruitType  
FROM Fruits1
```

UNION

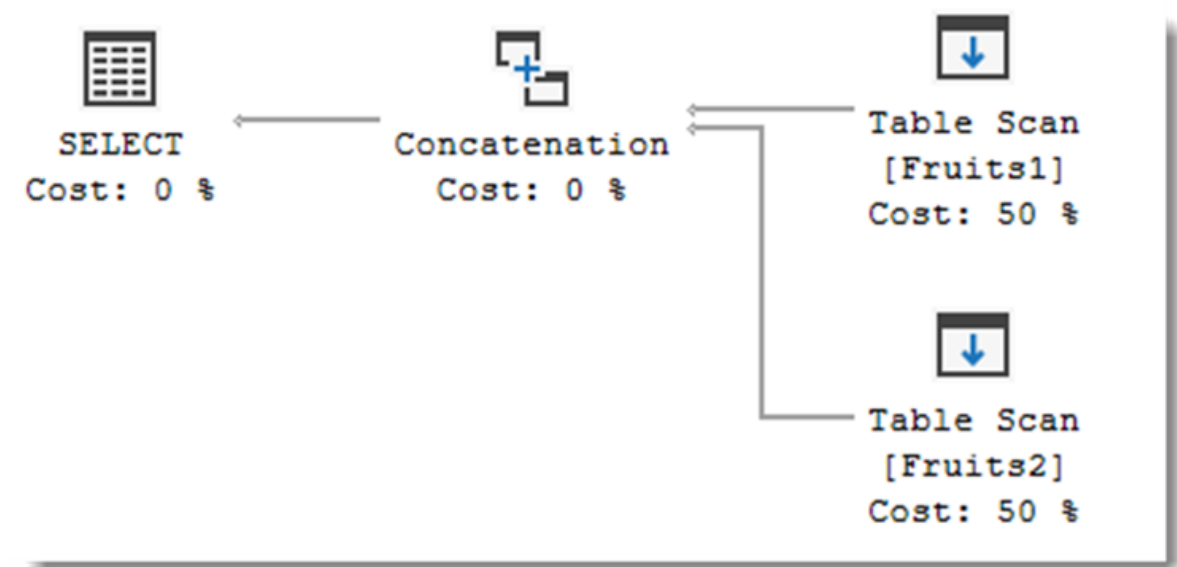
```
SELECT FruitName, FruitType  
FROM Fruits2;
```



```
SELECT FruitName, FruitType  
FROM Fruits1
```

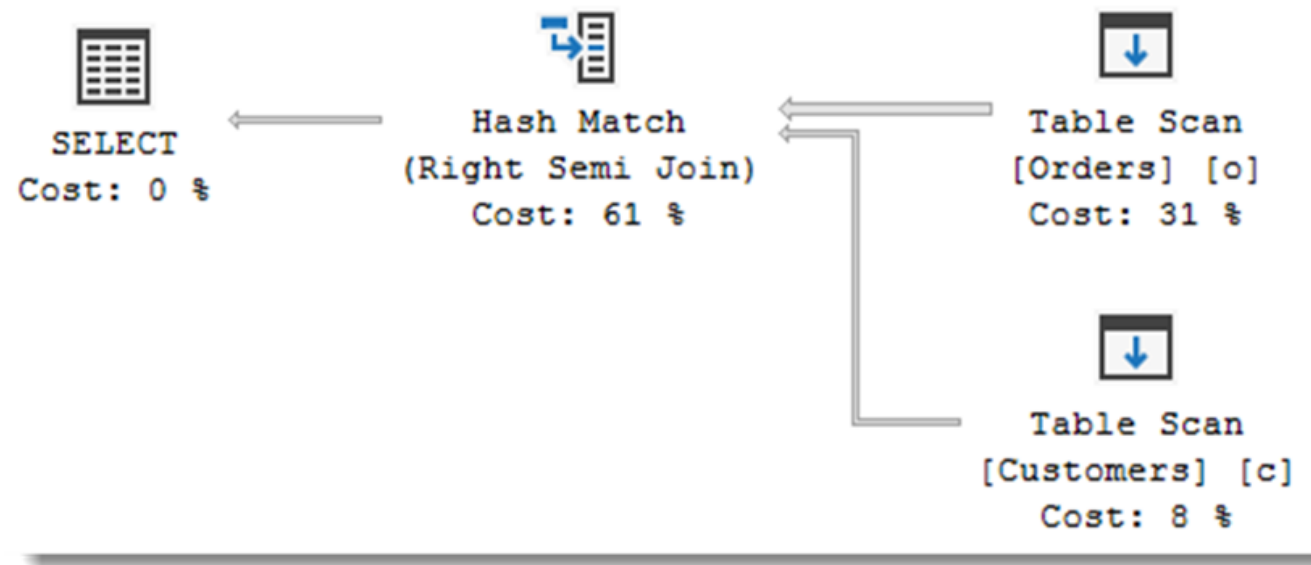
UNION ALL

```
SELECT FruitName, FruitType  
FROM Fruits2;
```

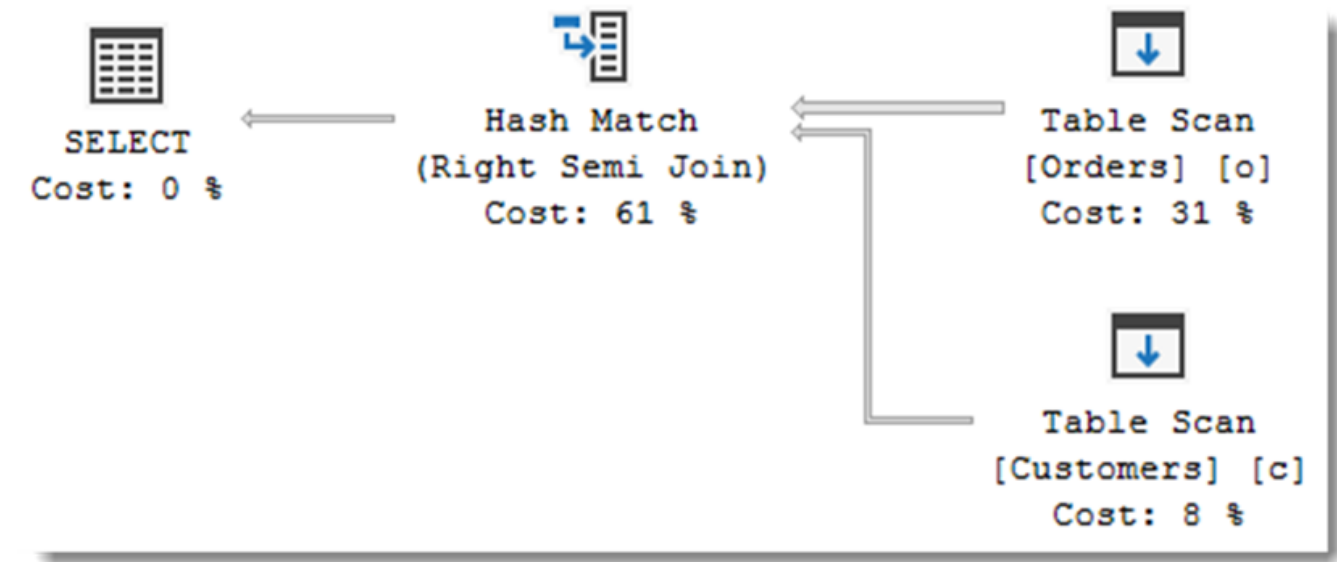


The same execution plan?

```
SELECT *  
FROM Customers  
WHERE CustomerID IN  
  (SELECT CustomerID  
   FROM Orders);
```



```
SELECT *  
FROM Customers c  
WHERE EXISTS  
  (SELECT 1  
   FROM Orders o  
   WHERE c.CustomerID = o.CustomerID);
```



Let's practice!

IMPROVING QUERY PERFORMANCE IN SQL SERVER

Query performance tuning: final notes

IMPROVING QUERY PERFORMANCE IN SQL SERVER

SQL

Dean Smith

Founder, Atamai Analytics

Final notes

SQL Server Execution Times:

CPU time = 390 ms, elapsed time = 382 ms.

- Time statistics examples and exercises presented in this chapter are reported in milliseconds.
- In the real world, it would not be uncommon to work with large complex queries that run for ten minutes, one hour or more.

Final notes

- Query statistics, indexes, and execution plans are all advanced topics with many books and websites devoted to each one.
- Communicate with your database administrator regarding the permissions required to use query performance tuning tools and commands in SQL Server.
- Don't rely on one tool or command for query performance tuning. They can often complement one another.

Conclusion

IMPROVING QUERY PERFORMANCE IN SQL SERVER